

Kalman Filters

Colin Shaw

04.27.2017

1. Introduction

- a. Welcome to Computer Science Club
- b. Thank you to RevUnit for sponsoring

2. What are Kalman filters used for?

- a. General time series analysis
 - i. Econometrics
- b. Motion planning
 - i. Guidance
 - ii. Navigation
 - iii. Apollo lander and nuclear missiles
- c. Sensor fusion
 - i. Lidar
 - ii. Radar

3. How Kalman filters fit in the larger world

- a. Measurement
 - i. Measuring the world around you
 - ii. Coordinates relative to you
 - iii. Examples
 - 1. Kalman filters
 - a. Plain (linear) Kalman filters
 - b. Extended Kalman filters
 - c. Unscented Kalman filters
 - 2. Double Exponential Smoothing
 - 3. Recursive Total Least Squares
- b. Localization
 - i. Measuring your position in the world
 - ii. Your position in larger map coordinates
 - iii. Examples
 - 1. Particle filters
 - 2. Histogram filters
- c. Both measurement *and* localization
 - i. Simultaneous Localization and Measurement (SLAM)

4. Torrid history of the introduction of the Kalman filter

- a. Rudolph Emil Kalman's 1960 paper
 - i. Hidden model was looked at with great skepticism
 - ii. Not met with open arms in the community
 - iii. EE journals would not publish it (ended up in ME)
- b. Unscented Kalman filters
 - i. Met with skepticism
 - ii. Again hard to publish

5. What type of process is a Kalman filter?

- a. Linear quadratic estimator
 - i. Linear system
 - ii. RMSE can be minimized by minimizing quadratic
- b. Bayesian estimator
 - i. Prior
 - ii. Distribution
 - iii. Posterior
 - iv. Conjugate priors
- c. Markov process
 - i. Future probabilities determined by most recent state

6. How are we presenting Kalman filters here?

- a. Developing intuition about a basic motion model
- b. Look at what we need to use the motion model
- c. Discuss noisy measurements
- d. Jump to talking about probability distributions
- e. Show the full set of Kalman filter equations
- f. Walk through a Kalman filter code example
- g. Discuss more complex motion models
- h. Discuss deviation from a linear model

7. Intuition regarding the motion model

- a. Draw measurements at time $t = 0, 1, 2$
- b. Infer measurement at time $t = 3$
- c. How do we know where the next point is?
 - i. Mental model is more than the measurements
 - ii. It has implied velocity
- d. How do we correct for errors
 - i. Error is the actual minus predicted (residual)
 - ii. Apply scaled residual to update model

8. More formality to the motion model

- a. Vector of position x and velocity v (or \dot{x})

$$\mathbf{x}_k = \begin{bmatrix} x \\ \dot{x} \end{bmatrix}$$

- b. State updates occur with a state transition function

$$\mathbf{F} = \begin{bmatrix} 1 & \Delta t \\ 0 & 1 \end{bmatrix}$$

- c. We can only observe measurables, so there is a measurement function \mathbf{H}

$$\mathbf{H} = \begin{bmatrix} 1 & 0 \end{bmatrix}$$

- d. The errors can be assumed to be *acceleration* noise
- There is a force causing the change in velocity
 - There is an acceleration associated with the force
 - This noise is distributed $N(0, \sigma_a^2)$

9. Adding the notion of noise

- a. What if we had noisy measurements only
- For each $t = 0, 1, 2, \dots$ we know some standard deviation of position
 - What does this give us?
 - The standard deviation at these positions
 - In effect, nothing
 - We need more information to be able to refine estimates of uncertainty
 - Think about the Bayesian estimation notion of the problem
 - We can use a prior and a new distribution to compute a posterior

10. Brief probability distribution review

- a. Gaussian.ipynb
- b. What is a normal distributed random variable
- Mean
 - Standard deviation
- c. Adding normally distributed random variables
- Analytical result
 - $\mu = \mu_1 + \mu_2$
 - $\sigma^2 = \sigma_1^2 + \sigma_2^2$
 - Pointwise addition
 - Convolution (ConvolutionAddition.ipynb)
- d. Multiplying Gaussians
- Analytical result
 - $\mu = \mu_1 \sigma_2^2 + \mu_2 \sigma_1^2 / \sigma_1^2 + \sigma_2^2$ (normalized weighted average)
 - $\sigma^2 = \sigma_1^2 \sigma_2^2 / \sigma_1^2 + \sigma_2^2$ (parallel sum)
 - PDF demonstration

11. Kalman filter model

a. Prediction

i. Predicted State Estimate

$$\hat{\mathbf{x}}_{k|k-1} = \mathbf{F}_k \hat{\mathbf{x}}_{k-1|k-1} + \mathbf{B}_k \mathbf{u}_k$$

ii. Predicted Estimate Covariance

$$\mathbf{P}_{k|k-1} = \mathbf{F}_k \mathbf{P}_{k-1|k-1} \mathbf{F}_k^T + \mathbf{Q}_k$$

b. Measurement

i. Innovation Residual

$$\tilde{\mathbf{y}}_k = \mathbf{z}_k - \mathbf{H}_k \hat{\mathbf{x}}_{k|k-1}$$

ii. Innovation Covariance

$$\mathbf{S}_k = \mathbf{H}_k \mathbf{P}_{k|k-1} \mathbf{H}_k^T + \mathbf{R}_k$$

iii. Optimal Kalman Gain

$$\mathbf{K}_k = \mathbf{P}_{k|k-1} \mathbf{H}_k^T \mathbf{S}_k^{-1}$$

iv. Updated State Estimate

$$\hat{\mathbf{x}}_{k|k} = \hat{\mathbf{x}}_{k|k-1} + \mathbf{K}_k \tilde{\mathbf{y}}_k$$

v. Updated Covariance Estimate

$$\mathbf{P}_{k|k} = (\mathbf{I} - \mathbf{K}_k \mathbf{H}_k) \mathbf{P}_{k|k-1}$$

12. Intuition about what the Kalman filter is doing

a. PositionVelocity.ipynb

b. Steps

- i. Initial position estimate (poor velocity knowledge)
- ii. Project to arbitrary velocity distribution
- iii. New measurement (decent position, poor velocity knowledge)
- iv. Posterior distribution now has much better position and velocity
- v. Repeat

c. Basic idea

- i. Additive noise (both prediction and measurement)
- ii. Multiplicative refinement in the Kalman gain calculation

13. Measuring error

- a. Root Mean Squared Error (RMSE)
 - i. How to compute

$$\text{RMSE} = \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2}$$

- ii. Is not normalized
 - iii. Depends on knowing ground truth
- b. Normalized Innovation Squared (NIS)
 - i. ChiSquared.ipynb
 - ii. How to compute

$$\text{NIS}_k = \tilde{y}_k^T S_k^{-1} \tilde{y}_k$$

- iii. Is normalized
 - iv. Does not depend on ground truth

14. Kalman filter code example

- a. KalmanFilter.ipynb
- b. Numpy for matrix arithmetic
- c. Steps
 - i. Initialize state and covariance
 - ii. Compute state prediction
 - iii. Update with measurement
 - iv. Repeat
- d. Error estimation

15. More complex motion models

- a. Sensor fusion
 - i. Merge two (or more) types of measurements
 - ii. Get better results than any provide alone
- b. Example
 - i. Radar
 - 1. Radius (inaccurate)
 - 2. Radial velocity (accurate)
 - 3. Angle (inaccurate)
 - ii. Lidar
 - 1. Radius (accurate)
 - 2. Angle (accurate)
- c. Measurement and state transfer function in general could be nonlinear

16. Deviation from the linear Kalman filter model

- a. GaussianTransformation.ipynb
- b. Extended Kalman filters
 - i. Linearize the nonlinear mappings
 - ii. Jacobian evaluation instead of matrix-vector multiplication
 - iii. Can be expensive to compute
 - iv. Not so good results in difficult cases
 - 1. Very nonlinear functions
 - 2. Non-normal distributions
- c. Unscented Kalman filters
 - i. Sample sigma points
 - ii. Map sigma points
 - iii. Estimator of distribution without knowing specifics
 - iv. Relatively inexpensive to compute
 - v. More stable results in most cases